UNIVERSITY of WISCONSIN-MADISON
Computer Sciences Department

CS 537   Andrea C. Arpaci-Dusseau
Introduction to Operating Systems   Remzi H. Arpaci-Dusseau

VIRTUALIZING MEMORY: FASTER WITH TLBS

Questions answered in this lecture:

Review paging...

How can page translations be made faster?

What is the basic idea of a TLB (Translation Lookaside Buffer)?

What types of workloads perform well with TLBs?

How do TLBs interact with context-switches?

ANNOUNCEMENTS

- P1: Due last Saturday : Graded soon

- Late handin directory for unusual circumstances

- Project 2: Available now
  - Due two weeks from yesterday: Monday, Oct 5
  - Can work with project partner in your discussion section (unofficial)
  - Two parts:
    - Linux: Shell -- fork() and exec(), file redirection, history
    - Xv6: Scheduler – simplistic MLFQ
    - Two discussion videos again; watch early and often!
  - **Fill out form on course web page if you would like project partner assigned (5:35 Wed)**
  - **Communicate with your project partner!**

- Exam 1: Two weeks, Thu 10/1 7:15 – 9:15 in **Humanities Bldg, Room 3650**
  - Class time that day for review
  - Look at homeworks / simulations for sample questions
  - **Fill out form on course web if you have academic conflict and must take alternate exam : DEADLINE THURSDAY; Notify Friday**

- Reading for today: Chapter 19

REVIEW: PAGING

Assume 4 KB pages

0x0000	PT	
0x0800	PT	
0x1000		ptbr
	P1	
0x2000		
	P2	
0x3000		
	P2	
0x4000		
	P1	
0x5000		
	P1	
0x6000		
	P2	
0x7000		

| 1 | 5 | 4 | ... | P1 pagetable |
| 6 | 2 | 3 | ... | P2 pagetable |

Virtual	Physical
load 0x0000	load 0x0800
	load 0x6000
load 0x1444	load 0x0808
	load 0x2444
load 0x1444	load 0x0008
	load 0x5444

What do we need to know?
Location of page table in memory (ptbr)
Size of each page table entry (assume 8 bytes)

REVIEW: PAGING PROS AND CONS

Advantages
- No external fragmentation
  - don't need to find contiguous RAM
- All free pages are equivalent
  - Easy to manage, allocate, and free pages

Disadvantages
- Page tables are too big
  - Must have one entry for every page of address space
- Accessing page tables is too slow [today's focus]
  - Doubles number of memory references per instruction

TRANSLATION STEPS

H/W: for each mem reference:

(cheap) 1. extract **VPN** (virt page num) from **VA** (virt addr)
(cheap) 2. calculate addr of **PTE** (page table entry)
(expensive) 3. read **PTE** from memory
(cheap) 4. extract **PFN** (page frame num)
(cheap) 5. build **PA** (phys addr)
(expensive) 6. read contents of **PA** from memory into register

Which steps are expensive?

Which expensive step will we avoid in today's lecture?
3) Don't always have to read PTE from memory!

EXAMPLE:
ARRAY ITERATOR

```
int sum = 0;

for (i=0; i<N; i++){
    sum += a[i];

}
```

Assume 'a' starts at 0x3000

Ignore instruction fetches

Observation:
Repeatedly access same PTE because program repeatedly
accesses same virtual page

What virtual addresses?

load 0x3000

load 0x3004

load 0x3008

load 0x300C

...

What physical addresses?

load 0x100C
load 0x7000
load 0x100C
load 0x7004
load 0x100C
load 0x7008
load 0x100C
load 0x700C

Aside: What can you infer?
• ptbr: 0x1000; PTE 4 bytes each
• VPN 3 -> PPN 7

STRATEGY: CACHE PAGE TRANSLATIONS

CPU

Translation
Cache

RAM

PT

Some popular entries

memory interconnect

TLB: **T**ranslation **L**ookaside **B**uffer

(yes, a poor name!)

TLB ORGANIZATION

TLB Entry

| Tag (virtual page number) | Physical page number (page table entry) |

Various ways to organize a 16-entry TLB (artificially small)

A

A   B

Index

A   B   C   D

Four-way set associative

Set

Two-way set associative

A   B   C   D   E   L   M   N   O   P

Fully associative

Direct mapped

Lookup
• Calculate set (tag % num_sets)
• Search for tag within resulting set

TLB ASSOCIATIVITY TRADE-OFFS

Higher associativity
- \+ Better utilization, fewer collisions
- − Slower
- − More hardware

Lower associativity
- \+ Fast
- \+ Simple, less hardware
- − Greater chance of collisions

TLBs usually fully associative

ARRAY ITERATOR (W/ TLB)

```
int sum = 0;

for (i = 0; i < 2048; i++){
    sum += a[i];

}
```

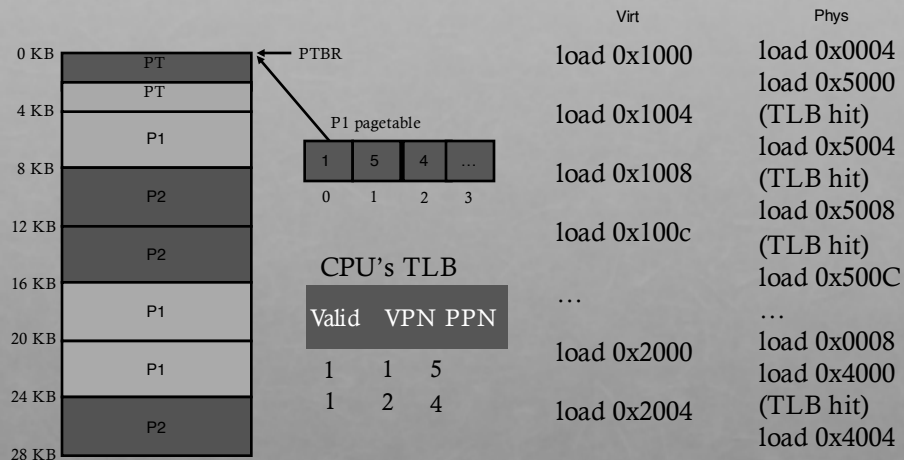Assume following virtual address stream:
load 0x1000

load 0x1004   What will TLB behavior look like?

load 0x1008

load 0x100C
…

TLB ACCESSES: SEQUENTIAL EXAMPLE

	Virt	Phys
0 KB — PT	load 0x1000	load 0x0004
PT		load 0x5000
4 KB — P1	load 0x1004	(TLB hit)
8 KB — P2	load 0x1008	load 0x5004 (TLB hit)
12 KB — P2	load 0x100c	load 0x5008 (TLB hit)
16 KB — P1	…	load 0x500C
20 KB — P1	load 0x2000	…
24 KB — P1	load 0x2004	load 0x0008 load 0x4000 (TLB hit)
28 KB — P2		load 0x4004

PTBR → PT

P1 pagetable

1	5	4	...
0	1	2	3

CPU's TLB

Valid	VPN	PPN
1	1	5
1	2	4

PERFORMANCE OF TLB?

```
int sum = 0;
for (i=0; i<2048; i++) {
      sum += a[i];
}
```

Calculate miss rate of TLB for data:
\# TLB misses / \# TLB lookups

\# TLB lookups?
= number of accesses to a = 2048

\# TLB misses?
= number of unique pages accessed
= 2048 / (elements of 'a' per 4K page)
= 2K / (4K / sizeof(int)) = 2K / 1K
= 2

Miss rate?
2/2048 = 0.1%

Hit rate? (1 – miss rate)
99.9%

Would hit rate get better or worse with smaller pages?
Worse

TLB PERFORMANCE

How can system improve TLB performance (hit rate) given fixed number of TLB entries?

Increase page size

Fewer unique page translations needed to access same amount of memory

TLB Reach:

Number of TLB entries * Page Size

TLB PERFORMANCE WITH WORKLOADS

Sequential array accesses almost always hit in TLB

- Very fast!

What access pattern will be slow?

- Highly random, with no repeat accesses

WORKLOAD
ACCESS PATTERNS

Workload A

```
int sum = 0;

for (i=0; i<2048; i++) {
     sum += a[i];

}
```

Workload B

```
int sum = 0;
srand(1234);
for (i=0; i<1000; i++) {
    sum += a[rand() % N];
}
srand(1234);
for (i=0; i<1000; i++) {
    sum += a[rand() % N];
}
```

WORKLOAD
ACCESS PATTERNS

Spatial Locality

Sequential Accesses

Temporal Locality

Repeated Random Accesses

WORKLOAD LOCALITY

Spatial Locality: future access will be to nearby addresses

Temporal Locality: future access will be repeats to the same data

What TLB characteristics are best for each type?

Spatial:
- Access same page repeatedly; need same vpn->ppn translation
- Same TLB entry re-used

Temporal:
- Access same address near in future
- Same TLB entry re-used in near future
- How near in future? How many TLB entries are there?

TLB
REPLACEMENT POLICIES

LRU: evict Least-Recently Used TLB slot when needed
(More on LRU later in policies next week)

Random: Evict randomly choosen entry

Which is better?

A	B	C	D	E	L	M	N	O	P

LRU TROUBLES

Valid	Virt	Phys
0	?	?
0	?	?
0	?	?
0	?	?

virtual addresses:

0 1 2 3 4

Workload repeatedly accesses same offset across 5 pages (strided access) but only 4 TLB entries

What will TLB contents be over time?
How will TLB perform?

TLB REPLACEMENT POLICIES

LRU: evict Least-Recently Used TLB slot when needed
  (More on LRU later in policies next week)

Random: Evict randomly choosen entry

Sometimes random is better than a "smart" policy!

TLB PERFORMANCE

How can system improve TLB performance (hit rate) given fixed number of TLB entries?

Increase page size
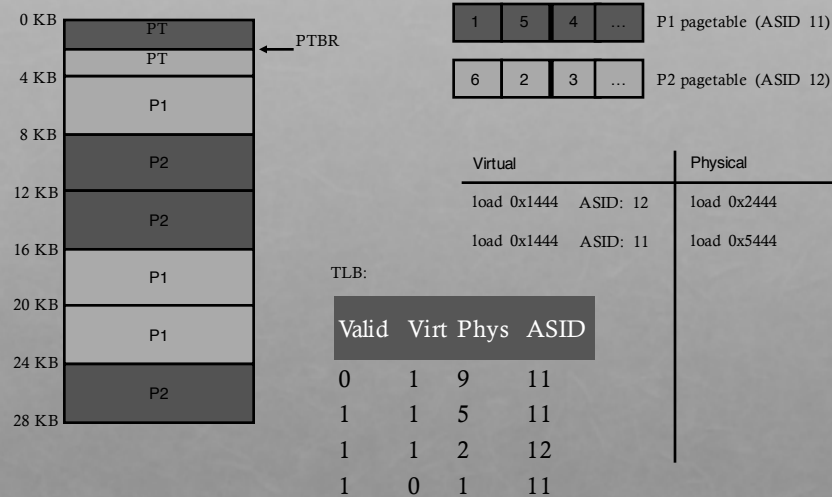 Fewer unique translations needed to access same amount of memory)

CONTEXT SWITCHES

What happens if a process uses cached TLB entries from another process?

Solutions?

1. Flush TLB on each switch
   - Costly; lose all recently cached translations

2. Track which entries are for which process
   - Address Space Identifier
   - Tag each TLB entry with an 8-bit ASID
     - how many ASIDs do we get?
     - why not use PIDs?

TLB EXAMPLE WITH ASID

0 KB	PT	← PTBR			
	PT				
4 KB	P1				
8 KB	P2				
12 KB	P2				
16 KB	P1				
20 KB	P1				
24 KB	P2				
28 KB					

1	5	4	...	P1 pagetable (ASID 11)
6	2	3	...	P2 pagetable (ASID 12)

Virtual		Physical
load 0x1444	ASID: 12	load 0x2444
load 0x1444	ASID: 11	load 0x5444

TLB:

Valid	Virt	Phys	ASID
0	1	9	11
1	1	5	11
1	1	2	12
1	0	1	11

TLB PERFORMANCE

Context switches are expensive

Even with ASID, other processes "pollute" TLB
- Discard process A's TLB entries for process B's entries

Architectures can have multiple TLBs
- 1 TLB for data, 1 TLB for instructions
- 1 TLB for regular pages, 1 TLB for "super pages"

HW AND OS ROLES

Who Handles TLB MISS? **H/W** or **OS**?

H/W: CPU must know where pagetables are
- CR3 register on x86
- Pagetable structure fixed and agreed upon between HW and OS
- HW "walks" the pagetable and fills TLB

OS: CPU traps into OS upon TLB miss
- "Software-managed TLB"
- OS interprets pagetables as it chooses
- Modifying TLB entries is privileged
  - otherwise what could process do?

Need same protection bits in TLB as pagetable
- rwx

SUMMARY

- Pages are great, but accessing page tables for every memory access is slow
- Cache recent page translations → TLB
  - Hardware performs TLB lookup on every memory access
- TLB performance depends strongly on workload
  - Sequential workloads perform well
  - Workloads with temporal locality can perform well
  - Increase **TLB reach** by increasing page size
- In different systems, hardware or OS handles TLB misses
- TLBs increase cost of context switches
  - Flush TLB on every context switch
  - Add ASID to every TLB entry

ANNOUNCEMENTS

- P1: Due last Saturday : Graded soon

- Late handin directory for unusual circumstances

- Project 2: Available now
  - Due two weeks from yesterday: Monday, Oct 5
  - Can work with project partner in your discussion section (unofficial)
  - Two parts:
    - Linux: Shell -- fork() and exec(), file redirection, history
    - Xv6: Scheduler – simplistic MLFQ
    - Two discussion videos again; watch early and often!
  - **Fill out form on course web page if you would like project partner assigned (5:35 Wed)**
  - **Communicate with your project partner!**

- Exam 1: Two weeks, Thu 10/1 7:15 – 9:15 in **Humanities Bldg, Room 3650**
  - Class time that day for review
  - Look at homeworks / simulations for sample questions
  - **Fill out form on course web if you have academic conflict and must take alternate exam : DEADLINE THURSDAY; Notify Friday**

- Reading for today: Chapter 19